

Machine Learning Lab 1

Альромхин Джорж, гр.858301

```
In [28]: 1 import numpy as np
          2 import pandas as pd
          3 import matplotlib.pyplot as plt
          4 from mpl_toolkits.mplot3d import Axes3D
          5 from matplotlib import cm
```

1. Load data

```
In [32]: 1 data_row1 = np.genfromtxt('../Data/lab1/ex1data1.txt', delimiter=',',
          2 data1 = pd.DataFrame(data_row1, columns=list(['Population', 'Profit']
          3 data1.head())
```

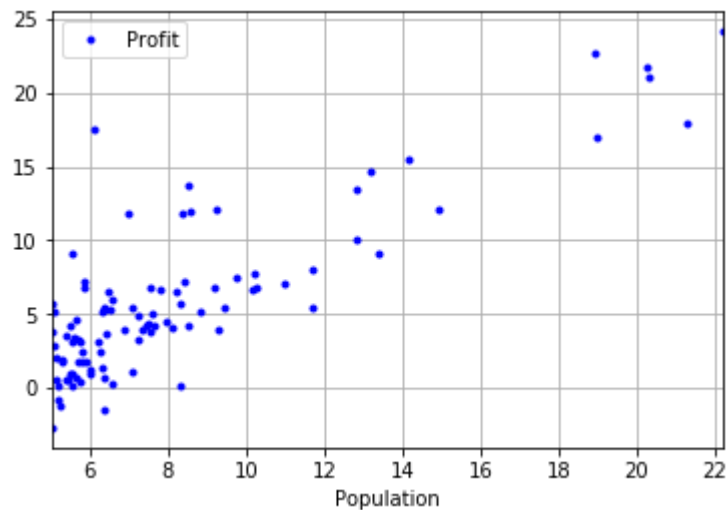
Out[32]:

	Population	Profit
0	6.1101	17.5920
1	5.5277	9.1302
2	8.5186	13.6620
3	7.0032	11.8540
4	5.8598	6.8233

2. Build a graph of the dependence of the profit of the restaurant on the population of the city in which it is located.

```
In [5]: 1 plt.figure()
2 data1.plot(x='Population', y='Profit', style=['b.'])
3 plt.grid(True)
4 plt.show()
```

<Figure size 432x288 with 0 Axes>



3. Cost function

```
In [12]: 1 def compute_cost(X, y, theta):
2         h = [np.matmul(x, theta.T).sum() for x in X]
3         return np.power(h - y, 2).sum() / (2 * m)
```

```
In [10]: 1 m = data_row1.shape[0] # Size of training set
2 n = data_row1.shape[1] # Size of feature vector
3 X1 = data1[['Population']]
4 X1.insert(0, 'theta_0', 1)
5 y1 = data1['Profit']
6 # theta the coefficients of the linear equation
7 theta = np.zeros((1, n))
8 X1.head()
```

Out[10]:

	theta_0	Population
0	1	6.1101
1	1	5.5277
2	1	8.5186
3	1	7.0032
4	1	5.8598

```
In [13]: 1 cost = compute_cost(X1.to_numpy(), y1.to_numpy(), theta)
2 print('Cost =', cost)
```

Cost = 32.072733877455676

4. Implement the gradient descent function to select model parameters. Build the resulting model (function) together with the graph

```
In [83]: 1 '''
2 function [theta, J_history] = gradient_dscent(X, y, theta, alpha, num
3 taking num_iters gradient steps with learning rate alpha
4 '''
5 def gradient_descent(X, y, theta, alpha, number_iterations):
6     m = y.shape[0]
7     n = X.shape[1]
8     j_history = []
9     for i in range(0, number_iterations):
10         deltas = np.zeros(n)
11         for j in range(0, n):
12             xj = X[:, j]
13             h = [np.matmul(x, theta.T)[0] for x in X]
14             deltas[j] = ((h - y) * xj).sum() * alpha / m
15             theta[0] -= deltas
16             j_history.append(compute_cost(X, y, theta))
17
18     return theta, j_history
```

```
In [17]: 1 iterations = 1500
2 alpha = 0.01
3 (theta, j_history) = gradient_descent(X1.to_numpy(), y1.to_numpy(),
```

```
In [18]: 1 print('gradient descent theras: ', theta)

gradient descent theras:  [[-3.63029144  1.16636235]]
```


```
In [53]: 1 print('for population = 35,000, profit prediction: ', (np.matmul([1,
2 print('for population = 70,000, profit prediction', (np.matmul([1, 7

for population = 35,000, profit prediction: 4519.767867701772
for population = 70,000, profit prediction 45342.45012944714
```

```
In [54]: 1 h = [np.matmul(x, theta.T).sum() for x in X1.to_numpy()]
2 data1_plot = data1.join(pd.DataFrame({'Linear regression': h}))
```

```
In [30]: 1 plt.figure()
2 ax = data1_plot.plot(x='Population', y='Profit', style=['b.'])
3 data1_plot.plot(x='Population', y='Linear regression', ax=ax)
4 plt.grid(True)
5 plt.show()
```

<Figure size 432x288 with 0 Axes>



5. Build a three-dimensional graph of the dependence of the loss function on the parameters of the model (θ_0 and θ_1) as a surface and as contours (contour plot).


```
In [36]: 1 theta0_vals = np.linspace(-10, 10, num=100)
2 theta1_vals = np.linspace(-1, 4, num=100)
3 # initialize J values to a matrix of 0
4 J_vals = np.zeros((theta0_vals.size, theta1_vals.size))
```

```
In [37]: 1 # Fill J values
2 for i in range(0, theta0_vals.size):
3     for j in range(0, theta1_vals.size):
4         t = np.array([[theta0_vals[i], theta1_vals[j]]])
5         J_vals[i, j] = computeCost(X1.to_numpy(), y1.to_numpy(), t)
6
7 J_vals = J_vals.T
```


```

In [40]: 1 fig = plt.figure()
          2 ax = fig.add_subplot(111, projection='3d')
          3 surf=ax.plot_surface(theta0_vals,theta1_vals,J_vals,cmap="coolwarm")
          4 fig.colorbar(surf, shrink=0.5, aspect=5)
          5 ax.set_xlabel(r'$\theta_0$')
          6 ax.set_ylabel(r'$\theta_1$')
          7 ax.set_zlabel(r'$J(\theta)$')
          8 ax.view_init(30,120)#rotate for better angle
          9 plt.show()

```



```
In [55]: 1 # Contour plot
2 plt.figure()
3 # Plot J_vals as 15 contours spaced logarithmically between 0.01 and
4 ax = plt.contour(theta0_vals, theta1_vals, J_vals, np.logspace(-2, 3
5 #plt.clabel(ax, inline=0, fontsize=2)
6 plt.xlabel(r'$\theta_0$')
7 plt.ylabel(r'$\theta_1$')
8 plt.plot(theta[0, 0], theta[0, 1], 'rx', linewidth=1, markersize=15)
9 plt.show()
```



6. Load Data

```
In [76]: 1 data_row2 = np.genfromtxt('../Data/lab1/ex1data2.txt', delimiter=',')
2 data2 = pd.DataFrame(data_row2, columns=list(['size', 'number', 'price']))
3 data2.head()
```

Out[76]:

	size	number	price
0	2104.0	3.0	399900.0
1	1600.0	3.0	329900.0
2	2400.0	3.0	369000.0
3	1416.0	2.0	232000.0
4	3000.0	4.0	539900.0

7. Normalize the signs. Did this affect the rate of convergence of the gradient descent? Give the answer in the form of a graph.

```
In [84]: 1 def normalization(X):
2     norm = (X - X.mean(axis=0)) / X.std(axis=0)
3     mu = X.mean(axis=0)
4     sigma = X.std(axis=0)
5     return norm, mu, sigma
6
```

```
In [79]: 1 X = data2[['size', 'number']]
2 X_norm, mu, sigma = normalization(X)
3 X_norm.describe()
```


```
Out[79]:
```

	size	number
count	4.700000e+01	4.700000e+01
mean	3.779483e-17	2.185013e-16
std	1.000000e+00	1.000000e+00
min	-1.445423e+00	-2.851859e+00
25%	-7.155897e-01	-2.236752e-01
50%	-1.417900e-01	-2.236752e-01
75%	3.376348e-01	1.090417e+00
max	3.117292e+00	2.404508e+00

```
In [80]: 1 y = data2['price']
2 m = y.size
3 n = data_row2.shape[1]
4 X.insert(0, 'theta_0', 1)
5 X_norm.insert(0, 'theta_0', 1)
6
7 theta1 = np.zeros((1, n))
8 theta2 = np.zeros((1, n))
```

```
In [81]: 1 (theta1, j_history) = gradient_descent(X.to_numpy(), y.to_numpy(), 1)
2 (theta2, j_norm_history) = gradient_descent(X_norm.to_numpy(), y.to_numpy(), 1)
```

```
In [82]: 1 p1 = plt.plot(range(0, len(j_history)), j_history, color='blue')
2 plt.legend('Raw')
3 p2 = plt.plot(range(0, len(j_norm_history)), j_norm_history, color='red')
4 plt.legend((p1[0], p2[0]), ('Raw', 'Normalized'))
5 plt.show()
```



8. Implement the loss function $J(\theta)$ and gradient descent for the case of multivariate linear regression using vectorization.

```
In [86]: 1 def gradient_descent2(X, y, theta, alpha, number_iterations):
2         m = y.shape[0]
3         j_history = []
4         XT = X.T
5         for i in range(0, number_iterations):
6             h = [np.matmul(x, theta.T)[0] for x in X]
7             loss = h - y
8             cost = np.sum(loss ** 2) / (2 * m)
9             gradient = np.matmul(XT, loss) / m
10            theta[0] -= alpha * gradient
11            j_history.append(cost)
12
13            return theta, j_history
```

```
In [87]: 1 iterations = 400
2 alpha = 0.01
3 theta_GD = np.zeros((1, n))
4
5 (theta_GD, j_history) = gradient_descent2(X_norm.to_numpy(), y.to_numpy())
6 print('Theta :', theta_GD)
```

Theta : [[334302.06399328 100087.11600585 3673.54845093]]

```
In [92]: 1 price = np.array([1, (1650 - mu[0]) / sigma[0], (4 - mu[1]) / sigma[1]])
2 print('price prediction for 4 rooms house: ', price)
```

price prediction for 4 rooms house: [294141.99995712]

9. Show that vectorization gives a performance boost

```
In [96]: 1 from timeit import default_timer as timer
2
3 iterations = 1000
4 alpha = 0.02
5 theta = np.zeros((1, n))
6
7 start = timer()
8 (theta, j_history) = gradient_descent(X_norm.to_numpy(), y.to_numpy())
9 end = timer()
10 exec_time1 = end - start
11 print('theta', theta)
12 print('time excuted:', exec_time1)
```

theta [[340412.65900156 110620.78816241 -6639.21215439]]
time excuted: 0.7502348749985686


```
In [95]: 1 theta = np.zeros((1, n))
2
3 start = timer()
4 (theta, j_history) = gradient_descent2(X_norm.to_numpy(), y.to_numpy)
5 end = timer()
6 exec_time2 = end-start
7
8 print('theta', theta)
9 print('time excuted:', exec_time2)
```


```
theta [[340412.65900156 110620.78816241 -6639.21215439]]
time excuted: 0.17354223399888724
```

```
In [105]: 1 diff = (exec_time1 / exec_time2)
2 print(diff, 'faster')
```

```
4.323067980117043 faster
```

10. Try changing the parameter α (learning factor). How does the graph of the loss function change depending on the number of iterations of the gradient descent? Draw the result as a graph.

```
In [101]: 1 alphas = np.linspace(0.1, 0.01, num=5)
2 plots = []
3 for alpha in alphas:
4     theta = np.zeros((1, n))
5     (theta, j_history) = gradient_descent2(X_norm.to_numpy(), y.to_r
6     p = plt.plot(range(0, len(j_history)), j_history)
7     plots.append(p[0])
8
9 plt.legend(plots, ["Alpha = %f" % (x) for x in alphas])
10 plt.show()
```



11. Build a model using an analytical solution that can be obtained by the least squares method. Compare the results of this model with the model obtained by gradient descent.

```
In [109]: 1 # computes the closed-form solution to linear regression using the normal equations
2 def normal_equations(X, y):
3     XX = np.asmatrix(X)
4     XT = XX.T
5     return ((XT @ XX).I @ XT) @ y
```

```
In [115]: 1 theta_A = normal_equations(X.to_numpy(), y.to_numpy())
2 print('theta from normal equations', theta_A)
3 print('theta from the normal normalized gradient descent: %s' % (theta_A))
4
5 price = np.array([1, 1200, 4]) @ theta_A.T
6 print('price prediction of 1200 sq. ft, 4 rooms house (using normal equations):', price)
7
8 price = np.array([1, (1200 - mu[0]) / sigma[0], (4 - mu[1]) / sigma[1]]) @ theta_A.T
9 print('price prediction of 1200 sq. ft, 4 rooms house (using gradient descent):', price)
```

theta from normal equations [[89597.9095428 139.21067402 -8738.01911233]]

theta from the normal normalized gradient descent: [[334302.06399328 100087.11600585 3673.54845093]]

price prediction of 1200 sq. ft, 4 rooms house (using normal equations): [[221698.64191464]]

price prediction of 1200 sq. ft, 4 rooms house (using gradient descent): [237467.6966757]