

Machine Learning Lab 2

Альромхин Джорж, гр.858301

```
In [65]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import scipy.optimize as opt
```

1. Load Data

```
In [26]: data_row1 = np.genfromtxt('../Data/lab2/ex2data1.txt', delimiter=',')
data1 = pd.DataFrame(data_row1, columns=list(['grade1', 'grade2', 'accepted']))
data1.head()
```


Out[26]:

	grade1	grade2	accepted
0	34.623660	78.024693	0.0
1	30.286711	43.894998	0.0
2	35.847409	72.902198	0.0
3	60.182599	86.308552	1.0
4	79.032736	75.344376	1.0

2. Build a graph where the axes are postponed grades in subjects, and the points are indicated by two different markers depending on whether the student entered the University or not.

```
In [29]: plt.figure()
ax = data1.loc[data1.accepted == 1].plot(x='grade1', y='grade2', style='b')
data1.loc[data1.accepted == 0].plot(x='grade1', y='grade2', style='r.')
plt.show()
```

<Figure size 432x288 with 0 Axes>



3. Implement loss functions $J(\theta)$ and gradient descent for logistic regression using vectorization.

```
In [30]: def sigmoid(X):
         return 1/(1+np.exp(-X))
```

```
In [31]: def cost_function(theta, X, y):
         m = y.size
         h = sigmoid(X @ theta)
         J = (-1 / m) * ((y.T @ np.log(h)) + ((1-y).T @ np.log(1-h)))
         error = h - y
         grad = (1 / m) * (X.T @ error)
         return J, grad
```

```
In [67]: m, n = data1.shape
X = data1[['grade1', 'grade2']]
X.insert(0, 'theta_0', 1)
X = X.to_numpy()
y = data1['accepted']
initial_theta = np.zeros(n)
cost, grad = cost_function(initial_theta, X, y)
print('cost at initial theta (zeros):', cost)
print('gradient at initial theta (zeros):', grad)
```

```
cost at initial theta (zeros): 0.6931471805599452
gradient at initial theta (zeros): [ -0.1         -12.00921659 -11.2628
 4221]
```

4. Implement other optimization methods (at least 2) for the implemented value function (e.g. Nelder-Mead Method, Broyden — Fletcher — Goldfarb — Shanno Algorithm, genetic

methods, etc.). It is allowed to use library implementations of optimization methods (for example, from the scipy library).

```
In [70]: def cost_optimize(theta, X, y):
         cost, _ = cost_function(theta, X, y)
         return cost

         def gradient_optimize(theta, X, y):
             _, grad = cost_function(theta, X, y)
             return grad

         def optimize(func, gradient, X, y, method):
             initial_theta = np.zeros(n)

             result = opt.minimize(fun=func, x0=initial_theta, args=(X, y), method=method)
             theta = result.x
             cost = func(theta, X, y)

             print(f'theta:\t{theta.ravel()}\ncost:\t{cost}')
             return result
```

Optimization using gradient information in a truncated Newton algorithm

```
In [71]: res = optimize(cost_optimize, gradient_optimize, X, y, 'TNC')
         res
```

```
theta: [-25.16131854  0.20623159  0.20147149]
cost:   0.20349770158947494
```

```
Out[71]:      fun: 0.20349770158947494
           jac: array([9.00038139e-09, 8.52696398e-08, 4.75764761e-07])
           message: 'Local minimum reached (|pg| ~= 0)'
           nfev: 36
           nit: 17
           status: 0
           success: True
           x: array([-25.16131854,  0.20623159,  0.20147149])
```

Unconstrained minimization of a function using the Newton-CG method

```
In [73]: result = optimize(cost_optimize, gradient_optimize, X, y, 'Newton-CG')
result

theta: [-25.16076473  0.20622705  0.20146712]
cost:  0.20349770163804756
```

```
Out[73]:      fun: 0.20349770163804756
      jac: array([-3.71475306e-06, -2.50437739e-04, -2.32431480e-04])
      message: 'Optimization terminated successfully.'
      nfev: 72
      nhev: 0
      nit: 29
      njev: 258
      status: 0
      success: True
      x: array([-25.16076473,  0.20622705,  0.20146712])
```

Implement the function of predicting the probability of a student's admission depending on the values of the exam scores.

```
In [74]: def predict(theta, X):
      predict = sigmoid(np.dot(X, theta))
      return predict > 0.5

theta = result.x
prob = sigmoid(np.array([1, 45, 85]) @ theta)
print(f'For a student with scores 45 and 85, we predict an admission probability of 0.7762868655447364')

For a student with scores 45 and 85, we predict an admission probability of 0.7762868655447364
```


```
In [75]: p = predict(np.array(theta), X)
print('Train Accuracy: %f' % ((y[p == y].size / float(y.size)) * 100.0))

Train Accuracy: 89.000000
```

Build a dividing line obtained as a result of training the model. Align the line with the graph from point 2.

```
In [77]: plt.figure()
ax = data1.loc[data1.accepted == 1].plot(x='grade1', y='grade2', style='b.')
data1.loc[data1.accepted == 0].plot(x='grade1', y='grade2', style='r.')
slope = -(theta[1] / theta[2])
intercept = -(theta[0] / theta[2])
x_vals = np.array(ax.get_xlim())
y_vals = intercept + (slope * x_vals)
plt.plot(x_vals, y_vals, c="k");
plt.show()
```

<Figure size 432x288 with 0 Axes>



Laod data2

```
In [79]: data_row2 = np.genfromtxt('../Data/lab2/ex2data2.txt', delimiter=',')
data2 = pd.DataFrame(data_row2, columns=list(['test1', 'test2', 'accepted']))
data2.head()
```


Out[79]:

	test1	test2	accepted
0	0.051267	0.69956	1.0
1	-0.092742	0.68494	1.0
2	-0.213710	0.69225	1.0
3	-0.375000	0.50219	1.0
4	-0.513250	0.46564	1.0

Build a graph where the test results are deferred along the axes, and the points are indicated by two different markers, depending on whether the product has passed the control or not.

```
In [80]: plt.figure()
ax = data2.loc[data2.accepted == 1].plot(x='test1', y='test2', style='r.', markersize=10)
data2.loc[data2.accepted == 0].plot(x='test1', y='test2', style='b.', markersize=10)
ax.set_ylabel('Test 2')
ax.set_xlabel('Test 1')
plt.show()
```

<Figure size 432x288 with 0 Axes>



Construct all possible combinations of features x_1 (the result of the first test) and x_2 (the result of the second test) in which the degree of the polynomial does not exceed 6, i.e. $1, x_1, x_2, x_1^2, x_2^2, x_1x_2, x_1^2x_2, x_1x_2^2, x_1^2x_2^2$, etc., (28 combinations in total).

```
In [85]: def map_feature(x1, x2):
    """
    Maps the two input features to quadratic features.
    Returns a new feature array with more features, comprising of
     $x_1, x_2, x_1^2, x_2^2, x_1x_2, x_1^2x_2, x_1x_2^2, x_1^2x_2^2$ , etc...
    Inputs  $x_1, x_2$  must be the same size
    """
    x1.shape = (x1.size, 1)
    x2.shape = (x2.size, 1)
    degree = 6
    out = np.ones(shape=(x1[:, 0].size, 1))

    m, n = out.shape

    for i in range(1, degree + 1):
        for j in range(i + 1):
            r = (x1 ** (i - j)) * (x2 ** j)
            out = np.append(out, r, axis=1)

    return out

mapped = map_feature(data2['test1'].to_numpy(), data2['test2'].to_numpy())
```

Implement L2 regularization for logistic regression and train it on an extended feature set by gradient descent.

```
In [113]: def cost_function_reg(theta, X, y, l):
    '''Compute the cost and partial derivatives as grads'''

    h = sigmoid(X.dot(theta))

    thetaR = theta[1:, 0]

    J = (-1.0 / m) * ((y.T @ np.log(h)) + ((1 - y.T) @ np.log(1.0 - h)))

    y.shape = h.shape
    delta = h - y
    sumdelta = delta.T @ X[:, 1]
    grad1 = (1.0 / m) * sumdelta

    XR = X[:, 1:X.shape[1]]
    sumdelta = delta.T @ XR

    grad = (1.0 / m) * (sumdelta + l * thetaR)

    out = np.zeros(shape=(grad.shape[0], grad.shape[1] + 1))

    out[:, 0] = grad1
    out[:, 1:] = grad

    return J.flatten(), out.T.flatten()

def gradient_descent(X, y, theta, l, alpha, num_iters):
    m = y.shape[0] # Size of training set
    j_history = []
    for i in range(0, num_iters):
        cost, grad = cost_function_reg(theta, X, y, l)
        grad.shape = theta.shape
        theta -= alpha * grad
        j_history.append(cost)

    return theta, j_history

m, n = data2.shape
y = data2['accepted'].to_numpy()
y.shape = (m, 1)
initial_theta = np.zeros(shape=(mapped.shape[1], 1))
#Set regularization parameter lambda to 1
l = 1
gd_theta, costs = gradient_descent(mapped, y, initial_theta, l, 0.2, 400)
```

Implement other optimization methods.

```
In [114]: def cost_decorator(theta, X, y):
            theta.shape = (mapped.shape[1], 1)
            cost, _ = cost_function_reg(theta, X, y, 1)
            theta.shape = (mapped.shape[1],)
            return cost

            def gradient_decorator(theta, X, y):
                theta.shape = (mapped.shape[1], 1)
                _, grad = cost_function_reg(theta, X, y, 1)
                theta.shape = (mapped.shape[1],)
                return grad

            initial_theta = np.zeros((mapped.shape[1], ))
            result = opt.minimize(fun=cost_decorator, x0=initial_theta, args=(mapped,))
            bfgs_theta = result.x
```

Implement the function of predicting the probability of passing the control of the product depending on the test results.

```
In [115]: p = predict(np.array(gd_theta), mapped)
            print('Train Accuracy with gradient descent: %f' % ((y[p == y].size / float(y.size)) * 100))
            p = predict(np.array(bfgs_theta), mapped)
            p.shape = y.shape
            print('Train Accuracy using the BFGS algorithm: %f' % ((y[p == y].size / float(y.size)) * 100))

            Train Accuracy with gradient descent: 80.508475
            Train Accuracy using the BFGS algorithm: 83.898305
```

Try different values of the regularization parameter λ . How does the choice of this value affect the appearance of the separating curve? Give the answer in the form of graphs.


```

In [117]: def train_with_plot(l):

    def plot(theta):
        u = np.linspace(-1, 1.5, 50)
        v = np.linspace(-1, 1.5, 50)
        z = np.zeros(shape=(len(u), len(v)))
        for i in range(len(u)):
            for j in range(len(v)):
                z[i, j] = (map_feature(np.array(u[i]), np.array(v[j]))).c

        z = z.T
        plt.figure()
        ax = data2.loc[data2.accepted == 1].plot(x='test1', y='test2', s=10)
        data2.loc[data2.accepted == 0].plot(x='test1', y='test2', style='o')
        ax.contour(u, v, z)
        ax.set_ylabel('Test 2')
        ax.set_xlabel('Test 1')
        plt.title(f'lambda = {l}')
        plt.show()


    initial_theta = np.zeros(shape=(mapped.shape[1], 1))
    theta, _ = gradient_descent(mapped, y, initial_theta, l, 0.2, 1000)
    data2 = pd.DataFrame(data_row2, columns=list(['test1', 'test2', 'acc

    plot(theta)

train_with_plot(0)


```

<Figure size 432x288 with 0 Axes>



```
In [118]: train_with_plot(100)
```

<Figure size 432x288 with 0 Axes>



load data3

```
In [106]: import scipy.io

data3 = scipy.io.loadmat('../Data/lab2/ex2data3.mat')
x = np.array(data3['X'])
y = np.squeeze(data3['y'])
np.place(y, y == 10, 0)
n = x.shape[1]
m = x.shape[0]
labels_count = 10
```

Visualize multiple random images from a dataset. The visualization must contain each digit at least once.

```
In [120]: import matplotlib.image as mpimg

subplots = 64
draw_seed = np.random.randint(low=0, high=x.shape[0], size=subplots)
draw_rows = x[draw_seed]
fig, ax = plt.subplots(8, 8, figsize=(8, 8))
for i, axi in enumerate(ax.flat):
    data = np.reshape(draw_rows[i], (20, 20), order='F')
    axi.imshow(data, cmap='binary')
    axi.set(xticks=[], yticks=[])

plt.show()
```



Implement a binary classifier using logistic regression using vectorization (loss function and gradient descent).

Add L2 regularization to the model.

Implement a one-versus-all multi-class classification. ¶

Implement the class prediction function on the image using trained classifiers.

```
In [ ]: X = np.ones(shape=(x.shape[0], x.shape[1] + 1))
X[:, 1:] = x
classifiers = np.zeros(shape=(labels_count, n + 1))
for i in range(0, labels_count):
    label = (y == i).astype(int)
    initial_theta = np.zeros(shape=(X.shape[1], 1))
    theta, costs = gradient_descent(X, label, initial_theta, 0.4, 2.8, 1)
    classifiers[i, :] = np.squeeze(theta)
```

```
In [ ]: def predict_class(input, classifiers):
    class_probs = sigmoid(input @ classifiers.transpose())
    if len(class_probs.shape) == 1:
        class_probs.shape = (1, class_probs.shape[0])
    predictions = class_probs.argmax(axis=1)
    return predictions
```

The percentage of correct classifications in the training sample should be about 95%.

```
In [229]: predictions = predict_class(X, classifiers)
print(f'Training accuracy: {str(100 * np.mean(predictions == y))}')

Training accuracy: 94.74000000000001
```